


Parallel DBMS

Chapter 22, Sections 22.1–22.6


Why Parallel Access To Data?

At 10 MB/s
1.2 days to scan

1,000 x parallel
1.5 minute to scan.



Parallelism:
divide a big problem
into many smaller ones
to be solved in parallel.




Parallel DBMS: Introduction


- ❖ Centralized DBMSs assume that
 - Processing of individual transactions is sequential
 - Control and/or data are maintained at one single site
- ❖ These assumptions have been relaxed in recent decades:
 - **Parallel DBMSs:**
 - ◆ Use of parallel evaluation techniques; parallelization of various operations such as data loading, index construction, and query evaluations.
 - ◆ Data may still be centralized; distribution dictated solely by performance considerations
 - **Distributed DBMSs:**
 - ◆ Use of both control and data distribution; data and control are dispersed and stored across several sites
 - ◆ Data distribution also dictated by considerations of increased availability and local ownership

Architectures for Parallel Databases


Shared Memory (SMP)



Shared Disk



Shared Nothing (network)



Easy to program
Expensive to build
Difficult to scaleup

Sequent, SGI, Sun

Hard to program
Cheap to build
Easy to scaleup


VMScluster, Sysplex

Tandem, Teradata, SP2

Some Parallelism Terminology


❖ Speed-Up

- For a given amount of data, more resources (CPUs) means proportionally more transactions processed per second.



❖ Scale-Up with DB size


- If resources increased in proportion to increase in data size, # of trans./sec. remains constant.



Some Parallelism Terminology

❖ Scale-Up

- If resources increased in proportion to increase in # of trans./sec., response time remains constant.




What Systems Work Which Way ?

(as of 9/1995)


Shared Nothing

Teradata:	400 nodes
Tandem:	110 nodes
IBM / SP2 / DB2:	128 nodes
Informix/SP2	48 nodes
ATT & Sybase	? nodes




Shared Disk

Oracle	170 nodes
DEC Rdb	24 nodes



Shared Memory

Informix	9 nodes
RedBrick	? nodes




Different Types of DBMS Parallelisms

- ❖ Intra-operator parallelism
 - get all machines working to compute a given operation (scan, sort, join)
- ❖ Inter-operator parallelism
 - each operator may run concurrently on a different site (exploits pipelining)
- ❖ Inter-query parallelism
 - different queries run on different sites
- ❖ We'll focus on intra-operator parallelism


Automatic Data Partitioning

Partitioning a table: Range




Good for equijoins,
range queries,
Selections, group-by

Hash



Good for equijoins

Round Robin



Good to spread load

**Shared disk and memory less sensitive to partitioning,
Shared nothing benefits from "good" partitioning**

Parallelizing Existing Code for Evaluating a Relational Operator

- ❖ How to readily parallelize existing code to enable sequential evaluation of a relational operator?
 - Idea: use parallel data streams
 - Details:
 - ◆ MERGE streams from different disks or the output of other operators to provide input streams for an operator
 - ◆ SPLIT output of an operator to parallelize subsequent processing
- ❖ A parallel evaluation plan is a dataflow network of relational, merge, and split operators.
 - Merge and split should have buffering capabilities
 - They should regulate the output of relational operators

Parallel Scanning and Bulk Loading

❖ Scanning

- Pages of a relation are read in parallel, and, if the relation is partitioned across many disks, retrieved tuples are merged.
- Selection of tuples matching some condition may not require all sites if range or hash partitioning is used.


❖ Bulk loading:

- Indexes can be built at each partition.
- Sorting of data entries required for building the indexes during bulk loading can be done in parallel.

Parallel Sorting

❖ Simple idea:

- Let each CPU sorts the portion of the relation located on its local disk;
- then, merge the sorted sets of tuples



❖ Better idea:

- First scan in parallel and redistribute the relation by range-partitioning all tuples; then each processor sorts its tuples:
 - ◆ The CPU collects tuples until memory is full
 - ◆ It sorts the in-memory tuples and writes out a run, until all incoming tuples have been written to sorted runs on disk
 - ◆ The runs on disk are then merged to get the sorted version of the portion of the relation assigned to the CPU
 - ◆ Retrieve the entire sorted relation by visiting the CPUs in the range-order to scan the tuples.
- Problem: how to avoid *skew* in range-partition!
- Solution: “sample” the data at start and sort the sample to determine partition points (splitting vector).

Parallel Joins: Range-Partition

- ❖ Assumptions: A and B are initially distributed across many processors, and k processors are available.
- ❖ Algorithm to join relations A and B on attribute age :
 1. At each processor where there are subsets of A and/or B , divide the range of age into k disjoint subranges and place partition A and B tuples into partitions corresponding to the subranges.
 2. Assign each partition to a processor to carry out a local join. Each processor joins the A and B tuples assigned to it.
 3. Tuples scanned from each processor are split, with a split operator, into output streams, depending on how many processors are available for parallel joins.
 4. Each join process receives input streams of A and B tuples from several processors, with merge operators merging all inputs streams from A and B , respectively.


Parallel Joins: Hash-Partitions

- ❖ Algorithm for hash-partition: Step 1 on previous slide changes to:
 1. At each processor where there are subsets of A and/or B , all local tuples are retrieved and hashed on the *age* attribute into k disjoint partitions, with the same hash function h used at all sites.
- ❖ Using range partitioning leads to a parallel version of Sort-Merge Join.
- ❖ Using hash partitioning leads to a parallel version of Hash Join.

Parallel Hash Join


Phase 1

Original Relations
(R then S)



- ❖ In first phase, partitions get distributed to different sites:
 - A good hash function *automatically* distributes work evenly!
- ❖ Do second phase at each site.
- ❖ Almost always the winner for equi-join.

Dataflow Network for Parallel Join



- ❖ Good use of split/merge makes it easier to build parallel versions of sequential join code.

Improved Parallel Hash Join

❖ Assumptions:

- A and B are very large, which leads to the size of each partition still being too large, which in turns leads to high local cost for processing the “smaller” joins.
- k partitions, n processors and $k=n$.

❖ Idea: Execute all smaller joins of A_i and B_i , $i=1,\dots,k$, one after the other, with each join executed in parallel using all processors.


Improved Parallel Hash Join (Cont'd)

❖ Algorithm:

1. At each processor, apply hash function $h1$ to partition A and B into partitions $i=1,\dots,k$. Suppose $|A| < |B|$. Then choose k such sum of all k partitions of A fits into aggregate memory of all n processors.
2. For $i=1,\dots,k$, process join of i -th partitions of A and B by doing this at every site:
 1. Apply 2nd hash function $h2$ to all A_i tuples to determine where they should be joined send t to site $h2(t)$.
 2. Add in-coming A_i tuples to an in-memory hash table.
 3. Apply $h2$ to all B_i tuples to determine where they should be joined send t to site $h2(t)$.
 4. Probe the in-memory hash table with in-coming B_i tuples.

Parallel Query Optimization

- ❖ Complex Queries: Inter-Operator parallelism
 - Pipelining between operators:
 - ◆ note that sort and phase 1 of hash-join block the pipeline!!
 - Bushy Trees



Parallel Query Optimization (Cont'd)

- ❖ Common approach: 2 phases
 - Pick best sequential plan (System R algorithm)
 - Pick degree of parallelism based on current system parameters.
- ❖ “Bind” operators to processors
 - Take query tree, “decorate” as in previous picture.

Parallel DBMS Summary

- ❖ Parallelism natural to query processing:
 - Both pipeline and partition parallelism!
- ❖ Shared-Nothing vs. Shared-Mem
 - Shared-disk too, but less standard
 - Shared-mem easy, costly. Doesn't scaleup.
 - Shared-nothing cheap, scales well, harder to implement.
- ❖ Intra-op, Inter-op, & Inter-query parallelism all possible.

Parallel DBMS Summary (Cont'd)

- ❖ Data layout choices important!
- ❖ Most DB operations can be done with partition-parallelism
 - Sort.
 - Sort-merge join, hash-join.
- ❖ Complex plans.
 - Allow for pipeline-parallelism, but sorts, hashes block the pipeline.
 - Partition parallelism achieved via bushy trees.

Parallel DBMS Summary (Cont'd)

- ❖ Hardest part of the equation: optimization.
 - 2-phase optimization simplest, but can be ineffective.
 - More complex schemes still at the research stage.
- ❖ We haven't said anything about transactions, logging.
 - Easy in shared-memory architecture.
 - Takes some care in shared-nothing.